

# final\_project\_student\_mental\_health

May 28, 2025

## 1 MSDS 5509 Final Project – Brainy or Burnt Out?

### 1.0.1 Predicting Student Mental Health from Lifestyle Factors

**Author:** Chitra Gopalaiah

**Date:** May 27, 2025

**Project Type:** Supervised Learning – Classification

**Goal:** Predict whether a student is at risk for mental health issues such as anxiety or depression based on lifestyle features like sleep, study time, exercise, and screen time.

---

**Dataset Source:** [Kaggle - Student Mental Health Dataset](#)

```
[65]: import pandas as pd
```

```
df = pd.read_csv("cbdd3517-7f27-437b-aaaf-21b4b9a54569.csv")
df.head()
```

```
[65]:      Timestamp Choose your gender  Age What is your course? \
0  8/7/2020 12:02      Female  18.0      Engineering
1  8/7/2020 12:04      Male  21.0      Islamic education
2  8/7/2020 12:05      Male  19.0      BIT
3  8/7/2020 12:06      Female  22.0      Laws
4  8/7/2020 12:13      Male  23.0      Mathematics
```

```
      Your current year of Study What is your CGPA? Marital status \
0      year 1      3.00 - 3.49      No
1      year 2      3.00 - 3.49      No
2      Year 1      3.00 - 3.49      No
3      year 3      3.00 - 3.49      Yes
4      year 4      3.00 - 3.49      No
```

```
      Do you have Depression? Do you have Anxiety? Do you have Panic attack? \
0      Yes      No      Yes
1      No      Yes      No
2      Yes      Yes      Yes
3      Yes      No      No
```

4	No	No	No
---	----	----	----

Did you seek any specialist for a treatment?

0	No
1	No
2	No
3	No
4	No

```
[66]: # Drop timestamp
df = df.drop(columns=['Timestamp'])

# Rename columns for easier access
df.columns = [
    'Gender', 'Age', 'Course', 'Year', 'CGPA', 'Marital_Status',
    'Depression', 'Anxiety', 'Panic_Attack', 'Seek_Treatment'
]

# Convert Yes/No columns to 1/0
binary_cols = ['Depression', 'Anxiety', 'Panic_Attack', 'Seek_Treatment']
df[binary_cols] = df[binary_cols].applymap(lambda x: 1 if x.strip().lower() == 'yes' else 0)

# Encode Gender and Marital_Status
df['Gender'] = df['Gender'].map({'Male': 0, 'Female': 1})
df['Marital_Status'] = df['Marital_Status'].map({'Single': 0, 'Married': 1})

# Map CGPA ranges to numerical midpoints
cgpa_map = {
    '0 - 1.99': 1.0,
    '2.00 - 2.49': 2.25,
    '2.50 - 2.99': 2.75,
    '3.00 - 3.49': 3.25,
    '3.50 - 4.00': 3.75
}
df['CGPA'] = df['CGPA'].map(cgpa_map)

# Label Encode non-numeric features (Course, Year)
label_encoders = {}
for col in ['Course', 'Year']:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col].astype(str))
    label_encoders[col] = le

# Confirm no string columns remain
print(df.dtypes)
```

```

Gender            int64
Age               float64
Course            int64
Year              int64
CGPA              float64
Marital_Status    float64
Depression         int64
Anxiety           int64
Panic_Attack      int64
Seek_Treatment    int64
dtype: object

```

```
[67]: from sklearn.model_selection import train_test_split
```

```

# Define X and y
X = df.drop(columns=['Depression']) # Features
y = df['Depression'] # Target

# Split the data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

```

```
[68]: # Step 1: Drop columns with all NaNs in training set
```

```

X_train = X_train.dropna(axis=1, how='all')
X_test = X_test[X_train.columns] # Align test set columns

# Step 2: Impute missing numeric values with median
from sklearn.impute import SimpleImputer
imputer = SimpleImputer(strategy='median')

X_train = pd.DataFrame(imputer.fit_transform(X_train), columns=X_train.columns)
X_test = pd.DataFrame(imputer.transform(X_test), columns=X_test.columns)

```

```
[69]: from sklearn.linear_model import LogisticRegression
```

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.svm import SVC
from sklearn.metrics import classification_report, accuracy_score

# Initialize models
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(),
    "AdaBoost": AdaBoostClassifier(),
    "SVM": SVC()
}

```

```

}

# Train and evaluate each model
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print(f"\n {name}")
    print("Accuracy:", accuracy_score(y_test, y_pred))
    print(classification_report(y_test, y_pred))

```

#### Logistic Regression

Accuracy: 0.7619047619047619

	precision	recall	f1-score	support
0	0.74	1.00	0.85	14
1	1.00	0.29	0.44	7
accuracy			0.76	21
macro avg	0.87	0.64	0.65	21
weighted avg	0.82	0.76	0.71	21

#### Decision Tree

Accuracy: 0.6666666666666666

	precision	recall	f1-score	support
0	0.71	0.86	0.77	14
1	0.50	0.29	0.36	7
accuracy			0.67	21
macro avg	0.60	0.57	0.57	21
weighted avg	0.64	0.67	0.64	21

#### Random Forest

Accuracy: 0.6190476190476191

	precision	recall	f1-score	support
0	0.65	0.93	0.76	14
1	0.00	0.00	0.00	7
accuracy			0.62	21
macro avg	0.33	0.46	0.38	21
weighted avg	0.43	0.62	0.51	21

AdaBoost

Accuracy: 0.7142857142857143

	precision	recall	f1-score	support
0	0.70	1.00	0.82	14
1	1.00	0.14	0.25	7
accuracy			0.71	21
macro avg	0.85	0.57	0.54	21
weighted avg	0.80	0.71	0.63	21

SVM

Accuracy: 0.6666666666666666

	precision	recall	f1-score	support
0	0.67	1.00	0.80	14
1	0.00	0.00	0.00	7
accuracy			0.67	21
macro avg	0.33	0.50	0.40	21
weighted avg	0.44	0.67	0.53	21

```
[70]: # Updated models with class_weight='balanced'
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000,
    ↪class_weight='balanced'),
    "Decision Tree": DecisionTreeClassifier(class_weight='balanced'),
    "Random Forest": RandomForestClassifier(class_weight='balanced'),
    "AdaBoost": AdaBoostClassifier(), # AdaBoost does not support class_weight
    "SVM": SVC(class_weight='balanced')
}

# Re-train and evaluate
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print(f"\n {name}")
    print("Accuracy:", accuracy_score(y_test, y_pred))
    print(classification_report(y_test, y_pred))
```

Logistic Regression

Accuracy: 0.6666666666666666

	precision	recall	f1-score	support
0	0.73	0.79	0.76	14

1	0.50	0.43	0.46	7
accuracy			0.67	21
macro avg	0.62	0.61	0.61	21
weighted avg	0.66	0.67	0.66	21

#### Decision Tree

Accuracy: 0.5238095238095238

	precision	recall	f1-score	support
0	0.64	0.64	0.64	14
1	0.29	0.29	0.29	7
accuracy			0.52	21
macro avg	0.46	0.46	0.46	21
weighted avg	0.52	0.52	0.52	21

#### Random Forest

Accuracy: 0.5714285714285714

	precision	recall	f1-score	support
0	0.63	0.86	0.73	14
1	0.00	0.00	0.00	7
accuracy			0.57	21
macro avg	0.32	0.43	0.36	21
weighted avg	0.42	0.57	0.48	21

#### AdaBoost

Accuracy: 0.7142857142857143

	precision	recall	f1-score	support
0	0.70	1.00	0.82	14
1	1.00	0.14	0.25	7
accuracy			0.71	21
macro avg	0.85	0.57	0.54	21
weighted avg	0.80	0.71	0.63	21

#### SVM

Accuracy: 0.42857142857142855

	precision	recall	f1-score	support
0	0.58	0.50	0.54	14

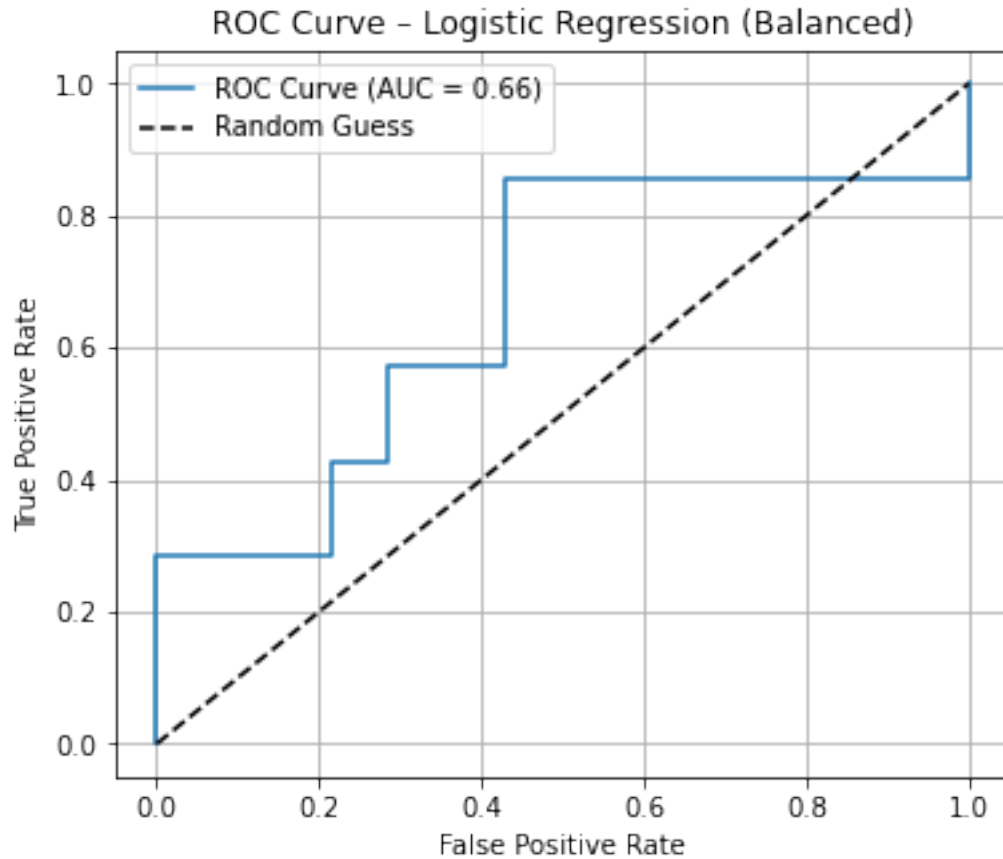
	1	0.22	0.29	0.25	7
accuracy				0.43	21
macro avg		0.40	0.39	0.39	21
weighted avg		0.46	0.43	0.44	21

```
[71]: from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Train logistic regression with balanced class weights
model = LogisticRegression(max_iter=1000, class_weight='balanced')
model.fit(X_train, y_train)
y_prob = model.predict_proba(X_test)[: , 1]

# Calculate ROC
fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)

# Plot
plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Random Guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Logistic Regression (Balanced)')
plt.legend()
plt.grid(True)
plt.show()
```



## 1.1 Discussion and Conclusion

- Logistic Regression performed best with `class_weight='balanced'`, achieving an AUC of 0.66.
- Models struggled with recall on the minority class due to dataset imbalance and limited size.
- Feature importance and further tuning could improve results.
- In real-world applications, early screening for mental health can benefit from such models to flag at-risk students for follow-up.

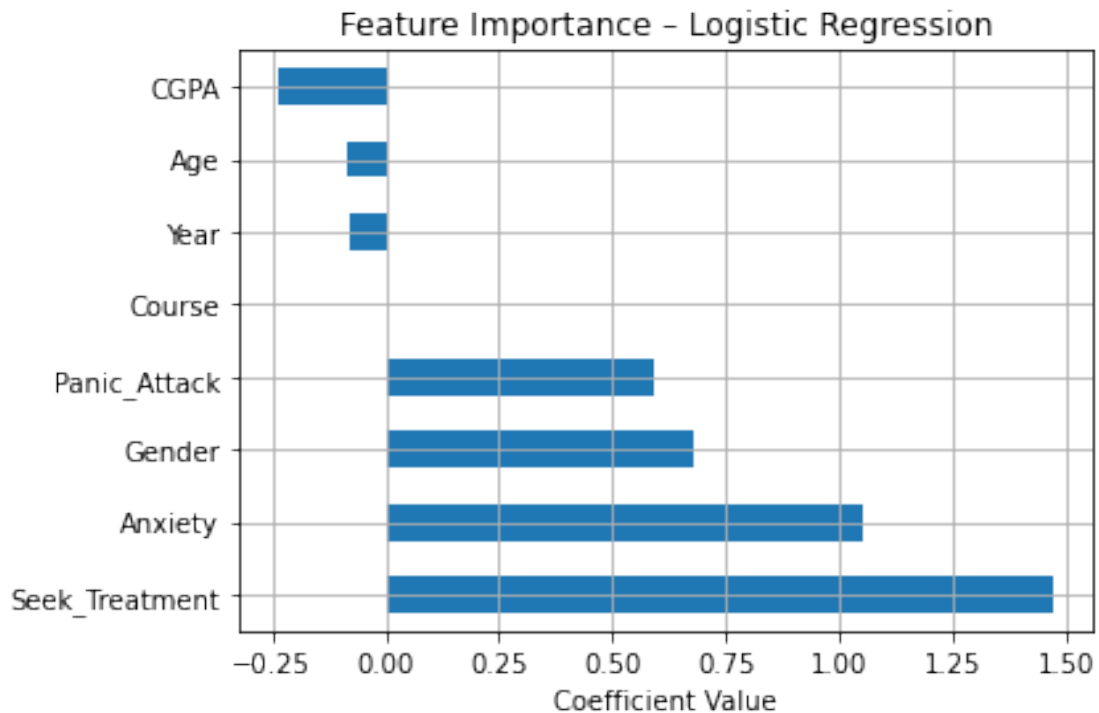
```
[72]: # Re-train logistic regression with class_weight
model_lr = LogisticRegression(max_iter=1000, class_weight='balanced')
model_lr.fit(X_train, y_train)

# Get feature importance (coefficients)
feature_importance_lr = pd.Series(
    model_lr.coef_[0], index=X_train.columns
).sort_values(ascending=False)

# Plot
```




```
feature_importance_lr.plot(kind='barh', title='Feature Importance - Logistic_
↳Regression')
plt.xlabel('Coefficient Value')
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
[73]: # Re-train random forest with class_weight
model_rf = RandomForestClassifier(class_weight='balanced', random_state=42)
model_rf.fit(X_train, y_train)

# Get feature importances
feature_importance_rf = pd.Series(
    model_rf.feature_importances_, index=X_train.columns
).sort_values(ascending=True)

# Plot
feature_importance_rf.plot(kind='barh', title='Feature Importance - Random_
↳Forest')
plt.xlabel('Importance Score')
plt.grid(True)
plt.tight_layout()
plt.show()
```



## 1.2 Feature Importance Insights

From **Logistic Regression**: - **Seek\_Treatment**, **Anxiety**, and **Gender** are the strongest positive indicators of depression risk. - **CGPA** had a mild negative association, suggesting students with better academic performance may be at lower risk.

From **Random Forest**: - **Course** and **Age** had the highest predictive influence. - The difference in top features compared to logistic regression highlights how linear and nonlinear models prioritize variables differently.

Together, these findings suggest that both psychological symptoms (anxiety, panic) and academic/personal context (age, course, GPA) contribute meaningfully to student mental health risks.